

(PSL) Coding Assignment 4

Contents

Part I: Gaussian Mixtures	1
Part II: HMM	4

Part I: Gaussian Mixtures

Objective

Implement the EM algorithm **from scratch** for a p -dimensional Gaussian mixture model with G components:

$$\sum_{k=1}^G p_k \cdot \mathcal{N}(x; \mu_k, \Sigma).$$

Requirements

Your implementation should consists of **four** functions.

- **Estep** function: This function should return an n -by- G matrix, where the (i, j) th entry represents the conditional probability $P(Z_i = k \mid x_i)$. Here i ranges from 1 to n and k ranges from 1 to G .
- **Mstep** function: This function should return the updated parameters for the Gaussian mixture model.
- **loglik** function: This function computes the log-likelihood of the data given the parameters.
- **myEM** function (main function): Inside this function, you can call the **Estep** and **Mstep** functions. The function should take the following inputs and return the estimated parameters and log-likelihood (via the **loglik** function):
 - **Input:**
 - * **data**: The dataset.
 - * **G**: The number of components. Although your code will be tested with $G = 2$ and $G = 3$, it should be able to handle any value of G . (You can, of course, ignore the case where $G > n$.)
 - * Initial parameters.
 - * **itmax**: The number of iterations.
 - **Output:**
 - * **prob**: A G -dimensional probability vector (p_1, \dots, p_G)
 - * **mean**: A p -by- G matrix with the k -th column being μ_k , the p -dimensional mean for the k -th Gaussian component.
 - * **Sigma**: A p -by- p covariance matrix Σ shared by all G components;

* **loglik**: A number equal to $\sum_{i=1}^n \log \left[\sum_{k=1}^G p_k \cdot \mathbf{N}(x; \mu_k, \Sigma) \right]$.

Implementation Guidelines:

The requirements are very similar to Coding Assignment 1. “No loops” means no explicit loops such as **for** or **while**, and no use of functions like **apply** or **map**.

- **Estep** function: No loops.
- **Mstep** function: You may only loop over **G** when updating **Sigma**.
- **loglik** function: You may only loop over **G**.
- You are not allowed to use pre-existing functions or packages for evaluating normal densities. However, you may use built-in functions to compute the inverse of a matrix or perform SVD.

Testing

Test your code with the provided dataset, [faithful.dat], with both $G = 2$ and $G = 3$.

For the case when $G = 2$, set your initial values as follows:

- $p_1 = 10/n$, $p_2 = 1 - p_1$.
- μ_1 = the mean of the first 10 samples; μ_2 = the mean of the remaining samples.
- Calculate Σ as

$$\frac{1}{n} \left[\sum_{i=1}^{10} (x_i - \mu_1)(x_i - \mu_1)^t + \sum_{i=11}^n (x_i - \mu_2)(x_i - \mu_2)^t \right].$$

Here $(x_i - \mu_i)$ is a 2-by-1 vector and the **superscript t denotes the transpose**. so the resulting Σ matrix is a 2-by-2 matrix.

Run your EM implementation with **20** iterations. Your results from **myEM** are expected to look like the following. (Even though the algorithm has not yet reached convergence, matching the expected results below serves as a validation that your code is functioning as intended.)

```
prob
[1] 0.04297883 0.95702117

mean
      [,1]      [,2]
eruptions 3.495642 3.48743
waiting   76.797892 70.63206

Sigma
      eruptions  waiting
eruptions 1.297936 13.92434
waiting   13.924336 182.58009

loglik
[1] -1289.569
```

For the case when $G = 3$, set your initial values as follows:

- $p_1 = 10/n$, $p_2 = 20/n$, $p_3 = 1 - p_1 - p_2$
- $\mu_1 = \frac{1}{10} \sum_{i=1}^{10} x_i$, the mean of the first 10 samples; $\mu_2 = \frac{1}{20} \sum_{i=11}^{30} x_i$, the mean of next 20 samples; and μ_3 = the mean of the remaining samples.
- Calculate Σ as

$$\frac{1}{n} \left[\sum_{i=1}^{10} (x_i - \mu_1)(x_i - \mu_1)^t + \sum_{i=11}^{30} (x_i - \mu_2)(x_i - \mu_2)^t + \sum_{i=31}^n (x_i - \mu_3)(x_i - \mu_3)^t \right].$$

Run your EM implementation with **20** iterations. Your results from **myEM** are expected to look like the following.

```

prob
[1] 0.04363422 0.07718656 0.87917922

mean
      [,1]      [,2]      [,3]
eruptions 3.510069 2.816167 3.545641
waiting   77.105638 63.357526 71.250848

Sigma
      eruptions  waiting
eruptions 1.260158 13.51154
waiting 13.511538 177.96419

loglik
[1] -1289.351

```

Derivation

Partial results for the derivation of the EM algorithm are given below. Note that the **faithful** data are two-dimensional, therefore $d = 2$, μ_k 's are 2-by-1 vectors and Σ is a 2-by-2 matrix.

1. The (marginal) likelihood function:

$$\begin{aligned}
& \prod_{i=1}^n p(x_i \mid p_{1:G}, \mu_{1:G}, \Sigma) \\
&= \prod_{i=1}^n [p_1 N(x_i; \mu_1, \Sigma) + \cdots + p_G N(x_i; \mu_G, \Sigma)] \\
&= \prod_{i=1}^n \left[p_1 \frac{\exp(-\frac{1}{2}(x_i - \mu_1)^t \Sigma^{-1} (x_i - \mu_1))}{\sqrt{(2\pi)^d |\Sigma|}} + \cdots + p_G \frac{\exp(-\frac{1}{2}(x_i - \mu_G)^t \Sigma^{-1} (x_i - \mu_G))}{\sqrt{(2\pi)^d |\Sigma|}} \right]
\end{aligned}$$

where $|\Sigma|$ denotes the determinant of matrix Σ . Your **loglik** function needs to compute the log of this function.

2. The complete likelihood function $\sum_{i=1}^n p(x_i, Z_i \mid p_{1:G}, \mu_{1:G}, \Sigma)$ or its log, which is the function we work with in the EM algorithm.

$$\begin{aligned}
& \prod_{i=1}^n p(x_i, Z_i \mid p_{1:G}, \mu_{1:G}, \Sigma) \\
&= \prod_{i=1}^n \prod_{k=1}^G \left[p_k \frac{\exp(-\frac{1}{2}(x_i - \mu_k)^t \Sigma^{-1} (x_i - \mu_k))}{\sqrt{(2\pi)^d |\Sigma|}} \right]^{1_{\{Z_i=k\}}}
\end{aligned}$$

3. Find the distribution of Z_i at the E-step. Given data and the current parameter value $(p_{1:G}^{(0)}, \mu_{1:G}^{(0)}, \Sigma^{(0)})$, Z_i follows a discrete distribute taking values from 1 to G with probabilities

$$\begin{aligned}
w_{ik} &:= P(Z_i = k \mid x_i, p_{1:G}^{(0)}, \mu_{1:G}^{(0)}, \Sigma^{(0)}) \\
&\propto P(x_i \mid Z_i = k, \mu_{1:G}^{(0)}, \Sigma^{(0)}) \times P(Z_i = k \mid p_{1:G}^{(0)})
\end{aligned}$$

4. The objective function you aim to maximize (or minimize) at the M-step. At the M-step, we optimize the following objective function (where the expectation is taken over Z_1, \dots, Z_n with respect to the probabilities computed at Step 3):

$$\begin{aligned}
 g(p_{1:G}, \mu_{1:G}, \Sigma) &= \mathbb{E} \log \prod_{i=1}^n p(x_i, Z_i \mid p_{1:G}, \mu_{1:G}, \Sigma) \\
 &= \mathbb{E} \sum_{i=1}^n \sum_{k=1}^G 1_{\{Z_i=k\}} \log \left[p_k \frac{\exp(-\frac{1}{2}(x_i - \mu_k)^t \Sigma^{-1} (x_i - \mu_k))}{\sqrt{(2\pi)^d |\Sigma|}} \right] \\
 &= \sum_{i=1}^n \sum_{k=1}^G w_{ik} \log \left[p_k \frac{\exp(-\frac{1}{2}(x_i - \mu_k)^t \Sigma^{-1} (x_i - \mu_k))}{\sqrt{(2\pi)^d |\Sigma|}} \right]
 \end{aligned}$$

where the last step is due to the fact that $\mathbb{E}[1_{\{Z_i=k\}}] = \mathbb{P}(Z_i = k) = w_{ik}$. You need to find the updating formulas for $p_{1:G}, \mu_{1:G}, \Sigma$ at the M-step.

Part II: HMM

Objective

Implement the Baum-Welch (i.e., EM) algorithm and the Viterbi algorithm **from scratch** for a Hidden Markov Model (HMM) that produces an outcome sequence of discrete random variables with three distinct values.

A quick review on parameters for Discrete HMM:

- **mx**: Count of distinct values X can take.
- **mz**: Count of distinct values Z can take.
- **w**: An mz-by-1 probability vector representing the initial distribution for Z_1 .
- **A**: The mz-by-mz transition probability matrix that models the progression from Z_t to Z_{t+1} .
- **B**: The mz-by-mx emission probability matrix, indicating how X is produced from Z .

Focus on updating the parameters **A** and **B** in your algorithm. The value for **mx** is given and you'll specify **mz**.

For **w**, initiate it uniformly but refrain from updating it within your code. The reason for this is that **w** denotes the distribution of Z_1 and we only have a single sample. It's analogous to estimating the likelihood of a coin toss resulting in heads by only tossing it once. Given the scant information and the minimal influence on the estimation of other parameters, we can skip updating it.

Baum-Welch Algorithm

The Baum-Welch Algorithm is the EM algorithm for the HMM. Create a function named **BW_onestep** designed to carry out the E-step and M-step. This function should then be called iteratively within **myBW**.

BW_onestep:

- **Input**:
 - data: a T-by-1 sequence of observations
 - Current parameter values
- **Output**:
 - Updated parameters: **A** and **B**

Please refer to formulas provided on Pages 7, 10, 14-16 in [lec_W7.2_HMM]

Viterbi Algorithm

This algorithm outputs the most likely latent sequence considering the data and the MLE of the parameters.

myViterbi:

- **Input:**
 - data: a T-by-1 sequence of observations
 - parameters: **mx**, **mz**, **w**, **A** and **B**
- **Output:**
 - **Z**: A T-by-1 sequence where each entry is a number ranging from 1 to **mz**.

Please refer to formulas provided on Pages 18-20 in [lec_W7.2_HMM]

Note on Calculations in Viterbi:

Many computations in HMM are based on the product of a sequence of probabilities, resulting in extremely small values. At times, these values are so small that software like R or Python might interpret them as zeros. This poses a challenge, especially for the Viterbi algorithm, where differentiating between magnitudes is crucial. If truncated to zero, making such distinctions becomes impossible. Therefore, it's advisable to evaluate these probabilities on a logarithmic scale in the Viterbi algorithm.

Testing

1. Test your code with the provided data sequence: [Coding4_part2_data.txt]. Set **mz** = 2 and start with the following initial values

$$w = \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix}, \quad A = \begin{pmatrix} 0.5 & 0.5 \\ 0.5 & 0.5 \end{pmatrix}, \quad B = \begin{pmatrix} 1/9 & 3/9 & 5/9 \\ 1/6 & 2/6 & 3/6 \end{pmatrix}$$

Run your implementation with **100** iterations. The results from your implementation of the Baum-Welch algorithm should match with the following:

A: the 2-by-2 transition matrix

```
0.49793938 0.50206062
0.44883431 0.55116569
```

B: the 2-by-3 emission matrix

```
0.22159897 0.20266127 0.57573976
0.34175148 0.17866665 0.47958186
```

The output from your Viterbi algorithm implementation should align with the following benchmarks. Please cross-check your results against the complete binary sequence available in [Coding4_part2_Z.txt]

```
1 1 1 1 1 1 1 2 1 1 1 1 2 2 1 1 1 1 1 1 2 2 2 2 1 1 1 1
1 1 2 1 1 1 1 1 1 1 2 2 1 1 1 1 1 2 2 2 1 1 1 1 2 2 2 2 1 1
.....
2 1 1 1 1 1 1 1
```

2. Initialize matrix B such that each entry is 1/3, and run your Baum-Welch algorithm for **20** and **100** iterations. Examine the resulting A and B matrices, and explain why you obtained these outcomes. Based on your findings, you should understand why we cannot initialize our parameters in a way that makes the latent states indistinguishable.